



# Serverless Computing – Cheap, Easy and Powerful

DIH-WEST

Sashko Ristov

# Agenda

- Introduction
  - About me
  - Introduction for the workshop
- FaaS
  - Task 1: Create a Simple Function
- BaaS Services
  - Task 2: Recognize faces from images
- Infrastructure-as-a-Code
  - Task 3: Deploy a function with Terraform
- Automation
  - Task 4: GitHub Actions + Terraform

# Positions



- (June 2022 - ) Assistant Professor (QE)
- 6 years Postdoctoral University Assistant (DPS)
- 3 years Assistant Professor (Skopje, N. Macedonia)
- 4 years Teaching and Research Assistant (Skopje, N. Macedonia)
- 8 years in Macedonian Telecom

# Publications in 2022

## • Journals / Magazines ×4

- **S. Ristov**, D. Kimovski, T. Fahringer, "FaaSinating Resilience for Serverless Function Choreographies in Federated Clouds," in *IEEE Trans. on Network and Service Management*, 2022.
- **S. Ristov**, S. Pedratscher, T. Fahringer, "xAFCL: Run scalable function choreographies across multiple FaaS systems," in *IEEE Transactions on Services Computing*, pp. 1–1, 2022
- S. Pedratscher, **S. Ristov**, T. Fahringer, "M2FaaS: Transparent and Fault Tolerant FaaSification of Node.js Monolith Code Blocks", in *Elsevier Fut. Gen. Comp. Sys.*, Accepted on Apr. 19, 2022.
- D. Kimovski, **S. Ristov**, R. Prodan, "Decentralized Machine Learning for Intelligent Health Care Systems on the Computing Continuum", in *IEEE Computer*, Accepted Jan. 2022.

## • Top Conferences ×3

- **S. Ristov**, M. Hautz, C. Hollaus, R. Prodan. "Simulate Serverless Workflows and Their Twins and Siblings in Federated FaaS," *ACM Symp. on Cloud Computing (SoCC'22)*, San Francisco, Nov. 2022
- **S. Ristov** and P. Gritsch. "FaaSSt: Optimize makespan of serverless workflows in federated commercial FaaS," (*IEEE CLUSTER '22*). Heidelberg, Germany, 182–194, 2022.
- **S. Ristov**, S. Brandacher, M. Felderer, R. Breu, "GoDeploy: Portable Deployment of Serverless Functions in Federated FaaS", *IEEE Cloud Summit 2022*, Fairfax, Virginia, USA, Oct. 2022

## • Workshops ×2

- **S. Ristov**, C. Hollaus, and M. Hautz. "Colder Than the Warm Start and Warmer Than the Cold Start! Experience the Spawn Start in FaaS Providers," in *ApPLIED '22*). Salerno, Italy, 35–39, 2022.
- M. Gusev, **S. Ristov**, et al. „CardioHPC: Serverless Approaches for Real-Time Heart Monitoring of Thousands of Patients“, *WORKS'22 Workshop (Supercomputing)* , Dallas, USA, Nov. 2022

# Serafin Plattner

- Bachelor student
- Bachelor thesis in serverless
- pyfOps: A pipeline for one-touch FaaSification, deployment, and testing of Python serverless functions
- On demand, we can show his pipeline, which integrates FaaSification and tests

# Organization (Social)

- Show agenda on WEB
- Refreshments
- Two Breaks

# Organization (Technical)

- FaaS familiarity?!
- GitHub account
- Fork the repository `https://github.com/FaaSTools/faas-tutorial`
- Install Terraform
- AWS Account

# Agenda

- Introduction
  - About me
  - Introduction for the workshop
- FaaS
  - Task 1: Create a Simple Function
- BaaS Services
  - Task 2: Recognize faces from images
- Infrastructure-as-a-Code
  - Task 3: Deploy a function with Terraform
- Automation
  - Task 4: GitHub Actions + Terraform



# FaaS Options



Alibaba



Google



AWS



IBM



# FaaS Options



AWS



# FaaS Options



AWS



# FaaS Options



AWS



# FaaS Options



AWS



AWS Step Functions

# FaaS Options



AWS



AWS Step Functions



S3 Storage

# Pay as you go

- **Upload** a function **in any language**
- NodeJS, Swift, Java, Go, Scala, Python, PHP, Ruby
- **Other** programming languages
  - Prepare your **own container**



# Charging?

- Fuel
  - Up to the 1/1000 (third digit after the "comma")





# Charges for AWS Lambda

## Requests

### 1M REQUESTS FREE

*First 1M requests per month are free.*

### \$0.20 PER 1M REQUESTS THEREAFTER

*\$0.0000002 per request.*

## Duration

### 400,000 GB-SECONDS PER MONTH FREE

*First 400,000 GB-seconds per month, up to 3.2M seconds of compute time, are free.*

### \$0.00001667 FOR EVERY GB-SECOND USED THEREAFTER

*The price depends on the amount of memory you allocate to your function.*

# Charges for AWS Lambda

## Requests

### 1M REQUESTS FREE

*First 1M requests per month are free.*

### **\$0.20 PER 1M REQUESTS THEREAFTER**

*\$0.0000002 per request.*

## Duration

### 400,000 GB-SECONDS PER MONTH FREE

*First 400,000 GB-seconds per month, up to 3.2M seconds of compute time, are free.*

### **\$0.00001667 FOR EVERY GB-SECOND USED THEREAFTER**

*The price depends on the amount of memory you allocate to your function.*

# Charges for AWS Lambda

## Requests

### 1M REQUESTS FREE

*First 1M requests per month are free.*

### **\$0.20 PER 1M REQUESTS THEREAFTER**

*\$0.0000002 per request.*

## Duration

### 400,000 GB-SECONDS PER MONTH FREE

*First 400,000 GB-seconds per month, up to 3.2M seconds of compute time, are free.*

### **\$0.00001667 FOR EVERY GB-SECOND USED THEREAFTER**

*The price depends on the amount of memory you allocate to your function.*

# Charges for AWS Lambda

| Requests  | Duration   |
|---|--|
| <b>1M REQUESTS FREE</b><br><i>First 1M requests per month are free.</i>     | <b>400,000 GB-SECONDS PER MONTH FREE</b><br><i>First 400,000 GB-seconds per month, up to 3.2M seconds of compute time, are free.</i>       |
| <b>\$0.20 PER 1M REQUESTS THEREAFTER</b><br><i>\$0.0000002 per request.</i> | <b>\$0.00001667 FOR EVERY GB-SECOND USED THEREAFTER</b><br><i>The price depends on the amount of memory you allocate to your function.</i> |

Free tiers both for requests and duration/resources

# Charges for AWS Lambda

| Requests  | Duration   |
|---|--|
| <b>1M REQUESTS FREE</b><br><i>First 1M requests per month are free.</i>     | <b>400,000 GB-SECONDS PER MONTH FREE</b><br><i>First 400,000 GB-seconds per month, up to 3.2M seconds of compute time, are free.</i>       |
| <b>\$0.20 PER 1M REQUESTS THEREAFTER</b><br><i>\$0.0000002 per request.</i> | <b>\$0.00001667 FOR EVERY GB-SECOND USED THEREAFTER</b><br><i>The price depends on the amount of memory you allocate to your function.</i> |

Free tiers both for requests and duration/resources

Some FaaS providers do not charge invocations!

# How to Code?

- Handler method
- Read keys from the JSON input data event ['key']
- Code the functions
- Return the output `returnJSON`
- **Show AWS Console**

# How to Deploy?

- Handler method format
- Region
- Memory
- Timeout
- Runtime Environment
- Role
- **Show AWS Console**

# How to Run?

- Test input
- Cold/Warm start
- Timeout
- Runtime Environment
- **Show AWS Console**



# Agenda

- Introduction
  - About me
  - Introduction for the workshop
- FaaS
  - Task 1: Create a Simple Function
- BaaS Services
  - Task 2: Recognize faces from images
- Infrastructure-as-a-Code
  - Task 3: Deploy a function with Terraform
- Automation
  - Task 4: GitHub Actions + Terraform

# How to Code?

- Boto3 for Python
- `boto3.resource('s3')`
- `boto3.client('rekognition')`
- layers
- Deployment package is ZIP

# Agenda

- Introduction
  - About me
  - Introduction for the workshop
- FaaS
  - Task 1: Create a Simple Function
- BaaS Services
  - Task 2: Recognize faces from images
- Infrastructure-as-a-Code
  - Task 3: Deploy a function with Terraform
- Automation
  - Task 4: GitHub Actions + Terraform

# Infrastructure as a Code

- Automate and manage
  - infrastructures
  - platforms
  - services

# Declarative vs. Imperative Approaches

- **Declarative** approach (**IDEMPOTENT**) - **Desired state** :
  - 20 servers
  - 2 FW
  - 3 VPCs
- **Imperative** approach - give **instructions**:
  - add 2 servers in AWS North Virginia
  - remove 2 servers in IBM Tokyo
  - add a FW
  - remove a security group

# Infrastructure Provisioning / Coding Tools / Standard

- Terraform
- Open Tosca
- AWS Cloud Formation (provider specific)
- AWS SAM (provider specific)
- Serverless.com

# Terraform

- tool for infrastructure provisioning, and platform and service deployment
- open source
- multi provider support

# Terraform

- tool for infrastructure provisioning, and platform and service deployment
- open source
- multi provider support
- **declarative** approach



# Terraform Use Cases

- **Provision** the environment
- **Update** the environment
- **Replicate** environment
  - From DEVELOPMENT to TESTING or PRODUCTION

# Terraform Architecture

- Inputs
  - **TF-config** (.tf)
    - What to create / configure
  - **TF-State**
- **Core** (how to reach the desired state from the current one?)
  - compares the config and the state
  - what needs to be deployed / updated / destroyed
- **Providers**
  - IaaS (AWS, IBM, ...)
  - PaaS (Kubernetes)
  - Serverless (FaaS)

# Terraform Actions

- `refresh`
  - query the provider about the current state
- create an execution plan (`init`)
  - which actions are necessary to be taken (no execution yet)
- `apply`
  - execute the plan
- `destroy`
  - destroy the infrastructure according to the plan (opposite from `apply`)

# Agenda

- Introduction
  - About me
  - Introduction for the workshop
- FaaS
  - Task 1: Create a Simple Function
- BaaS Services
  - Task 2: Recognize faces from images
- Infrastructure-as-a-Code
  - Task 3: Deploy a function with Terraform
- Automation
  - Task 4: GitHub Actions + Terraform

# GitHub Actions

- A pipeline to automatize actions after some push/pull on GitHub
- YAML-based script
- create `.github/workflows/main.yaml`
- <https://docs.github.com/en/actions/using-workflows/workflow-syntax-for-github-actions>

# GitHub Actions Example

Open Actions -> Python application

# QUESTIONS?



Email: [sashko.ristov@uibk.ac.at](mailto:sashko.ristov@uibk.ac.at)